



# Blockstore

**A data storage solution built with Blockchain**

Fatih Mehmet Bakır  
Kyle Carson  
CS 293S, UCSB Winter 2018

# Motivation

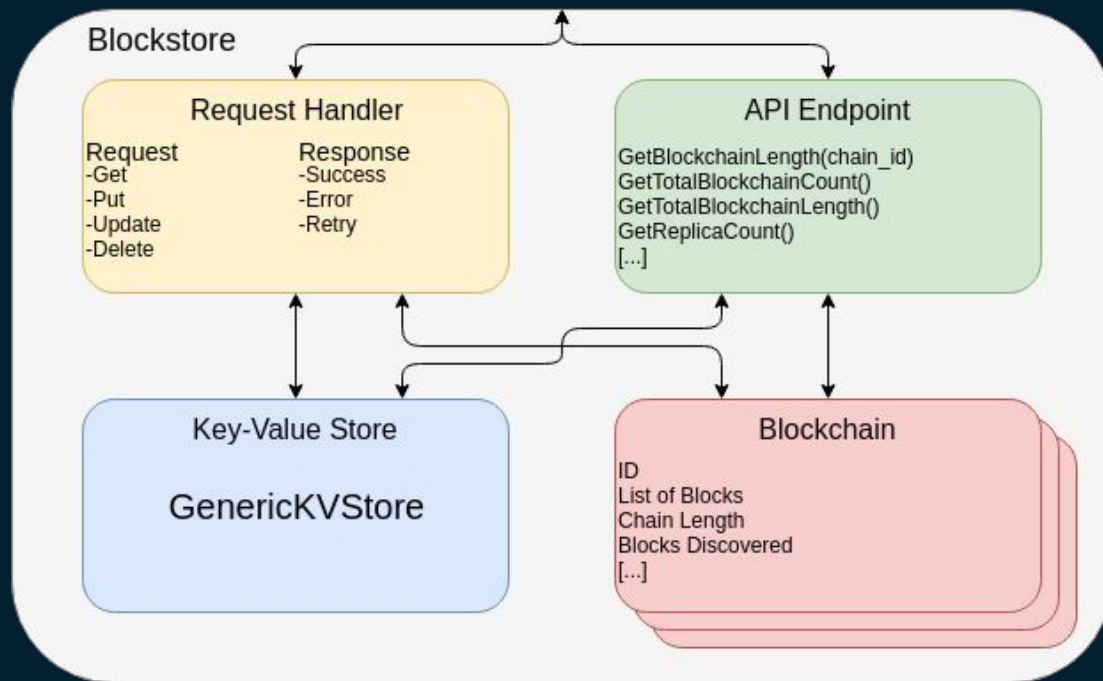
- › Consensus, replication, and sharding are exceedingly difficult in the context of data storage
  - › Blockchain takes an interesting approach to distributed communication/consensus
  - › A permissioned environment such as owning all of the mining/storage nodes provides some room for experimentation
    - › A semi-permissioned environment has the potential to cut costs
- ~~Blockchain is all the hype~~

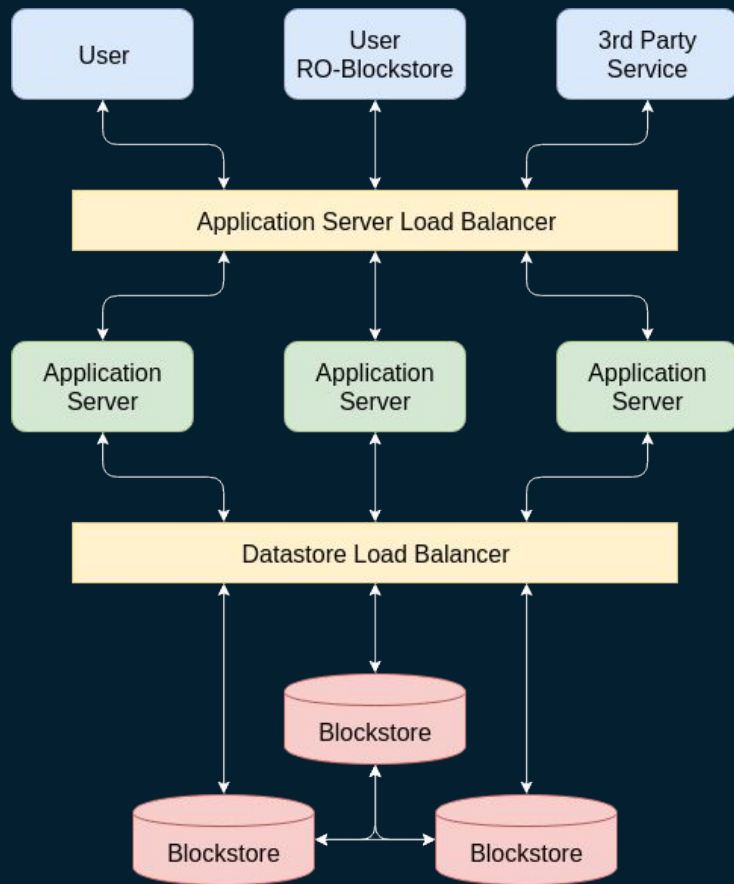
# What does Blockstore aim to achieve?

*Provide a simple key-value storage solution using a modified Blockchain implementation for managing operation logs and distributed communication*

# Overview

- › Modified Blockchain-variant using Generics
- › C++ for intensive operations such as mining, validating
- › NodeJS + TypeScript
  - › HTTP for client communication
  - › Socket.io for internal communication
- › Simple in-memory KV store using a dictionary
- › Request and API Handlers
- › Client module as an application-programmer's interface





# Typical Flow

- › Request is made by some client to Blockstore
- › Reads get sent directly to the KV-store, while any sort of Write is created in the blockchain as a unit of Operation
- › Operation added to a block and mined
- › Once mined, it is appended to the blockchain, distributed to all other known replicas, and a response is sent to the client
- › Mining is synchronous, rest is asynchronous

# Blockchain Modifications

- › No merckle-tree, need to maintain full history of operations
- › Significantly lower bounds on honest node count and mining difficulty requirements since there are no attackers
- › Full view/control of nodes in the network
- › Flexibility in the data size of Operations, Blocks, Async vs. Sync communication, etc via a cluster-wide configuration file



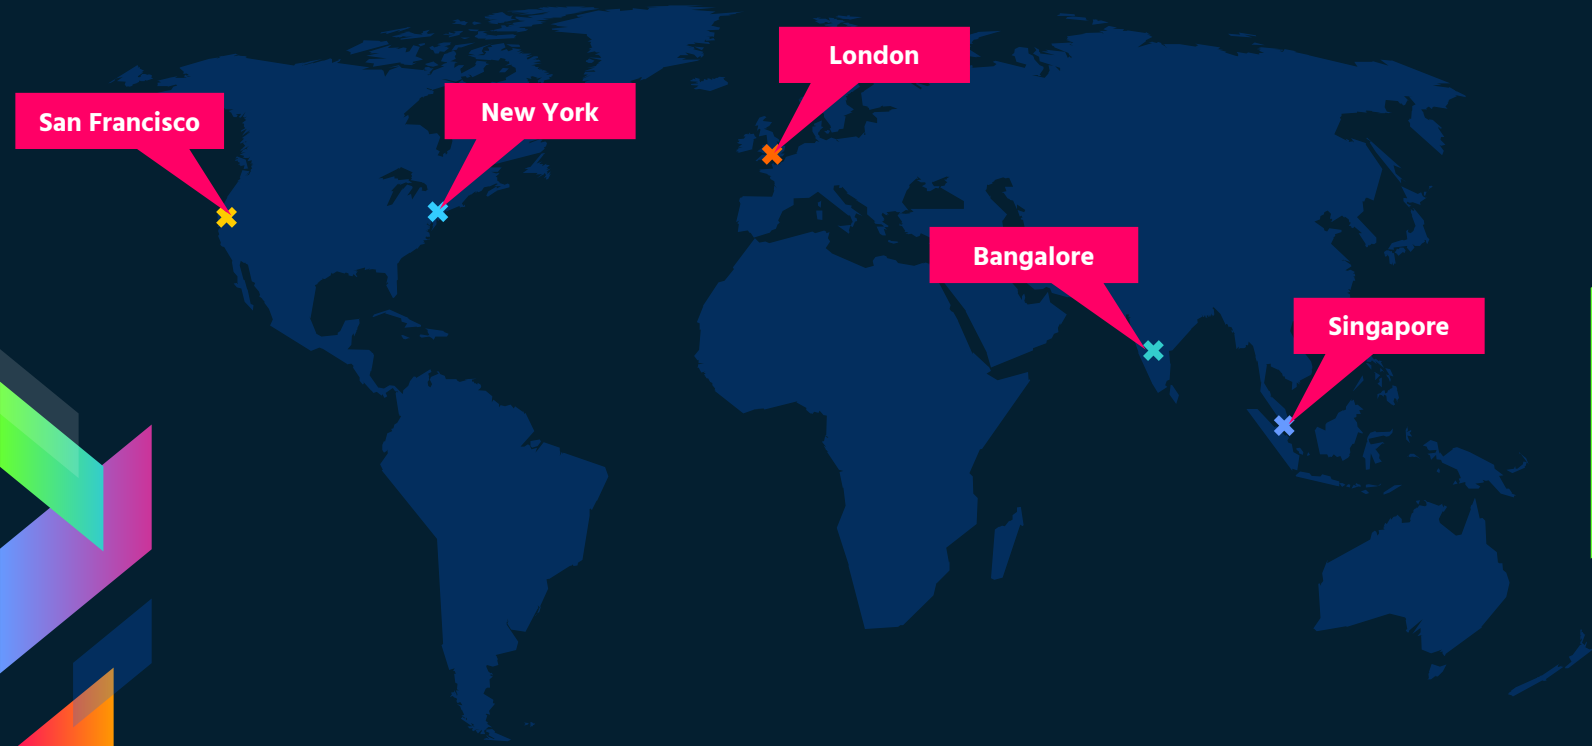
# Experimental Setup

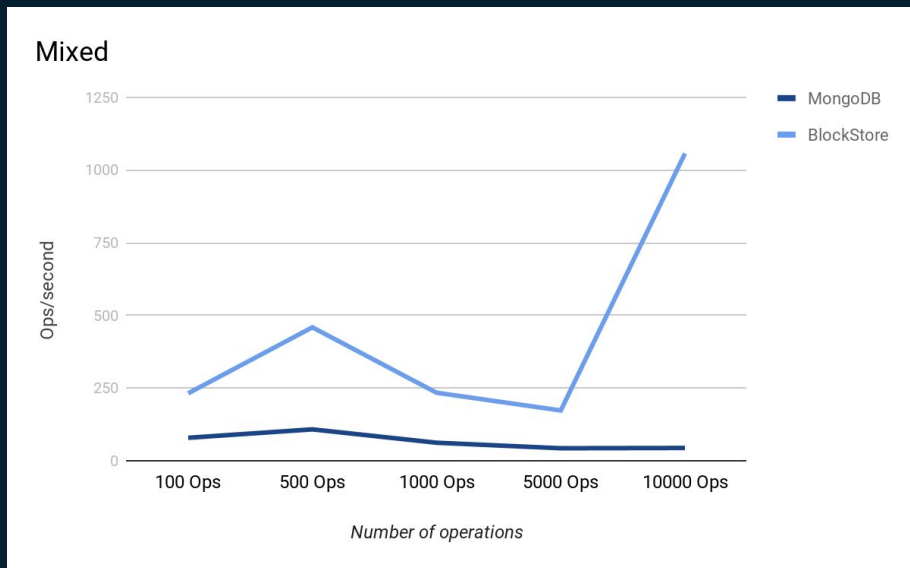
- › Remotely test MongoDB vs Blockstore with YCSB
  - › Latency
  - › Throughput
- › Multi-server cluster
- › Simple Demo Application

# Cluster Deployment


- › Docker images running on DigitalOcean
- › Servers located in multiple regions
- › VM Capacity:
  - › 1 CPU
  - › 1 GB RAM
  - › Ubuntu 16.04

# Geographical Deployment






**65% Get, 25% Put, 10% Update**



**90% Get, 5% Put, 5% Update**



**45% Put, 45% Update, 10% Get**

# Lessons Learned

- › Bitcoin is a very simplistic approach to Blockchain
- › Coupling of transaction commital to mining is *slow*
- › Strong consistency guarantees are difficult in a reasonable timeframe
  - › Decays into 2PC
- › Redundancy between contents of Blockchain and KV comes with overhead

# Additional Ideas to Explore

- › Read-only replica functionality
- › Garbage collecting the blockchain
  - › Checkpointing
- › Leveraging permissioned environment to do leader-based transaction committal (ByzCoin)
- › Build schema definition API for ease of development
- › Expand request query types to support more advanced queries on multiple items



<http://bs.fatihbakir.net>



**WE'LL DO IT LIVE!**

Abstract geometric shapes in the corners. The top-left and bottom-left corners feature overlapping translucent shapes in shades of green, blue, and orange. The top-right and bottom-right corners feature overlapping translucent shapes in shades of purple, blue, and orange.

# THANKS!

Questions?

@FatihBAKIR

@carsonkk

<https://github.com/FatihBAKIR/blockstore>